

PyCon Australia 2023

Building 3D Trusted Data Pipelines With Dagster, Dbt, and Duckdb

Danh Phan

Adelaide, August 2023

1

<https://github.com/danhphan/trusted-data-pipeline>

About me



Danh Phan

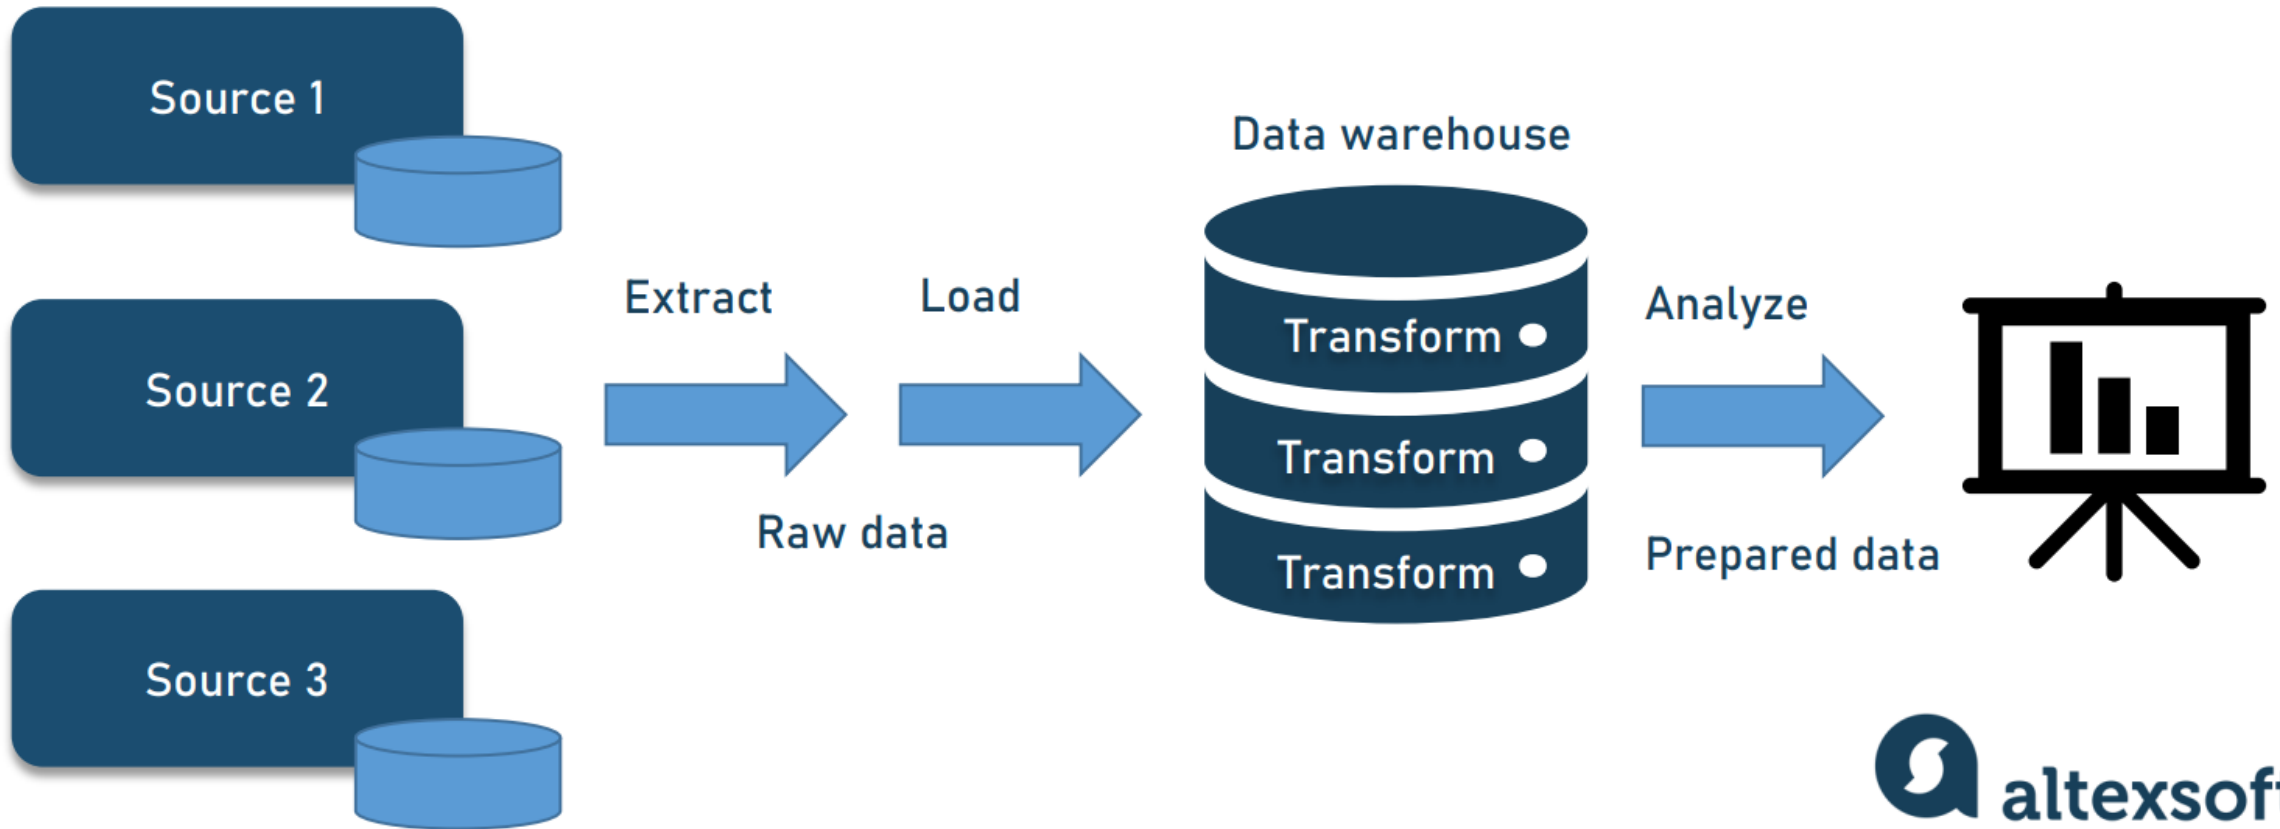
- Senior Data engineer at IDS, HESTA.
- An open-source contributor & a camper
- Ex. ML Researcher, Ex. Data analyst,
- Ex. Web dev, Ex. Database designer.
- <https://danhphan.net>

Views And Opinions Are My Own!

Overview

- Context
- Trusted data pipelines
- Demo

ELT PIPELINE





Need robust/reliable pipelines

Interesting talks on this topic

- 2022 Open Data Science, **Sam Bail** : [Building a Robust Data Pipeline with the "dag Stack": dbt, Airflow, and Great Expectations](#)
- 2021 Bigeye, **Egor Gryaznov** : [Data Reliability Engineering—Reliable Data Pipelines 101](#)
- 2021 DataEngBytes, **Harmeet Sokhi** : [Shift-left testing : Building reliable Data Pipelines](#)
- 2019 DataBricks, **Steven Yu** : [Building Robust Production Data Pipelines with Databricks Delta](#)
- 2017 DataBricks, **Xiao Li** : [Building Robust ETL Pipelines with Apache Spark](#)
- 2016 Jfokus, **Lars Albertsson** : [Data pipelines from zero to solid](#)
- ...

Lessons-learned

- Testing data pipelines (***)
- DevOps practice: code versioning, CI/CD
- Infrastructure as code (Terraform , CloudFormation, ...)
- Container environment (Docker)
- Data lineage and monitoring
- Enhance data contracts from upstream
- ...

Robust is not enough!

- Robust/reliable is not enough
- We absolutely need it
- But we also need to deliver the data quality to our data consumers (data analysts, BI developers, or Data scientists)

Trusted data pipelines also focus on data quality!

Robust pipeline vs. Trusted pipeline

Engineering focus



@istockphoto

Data consumers' focus



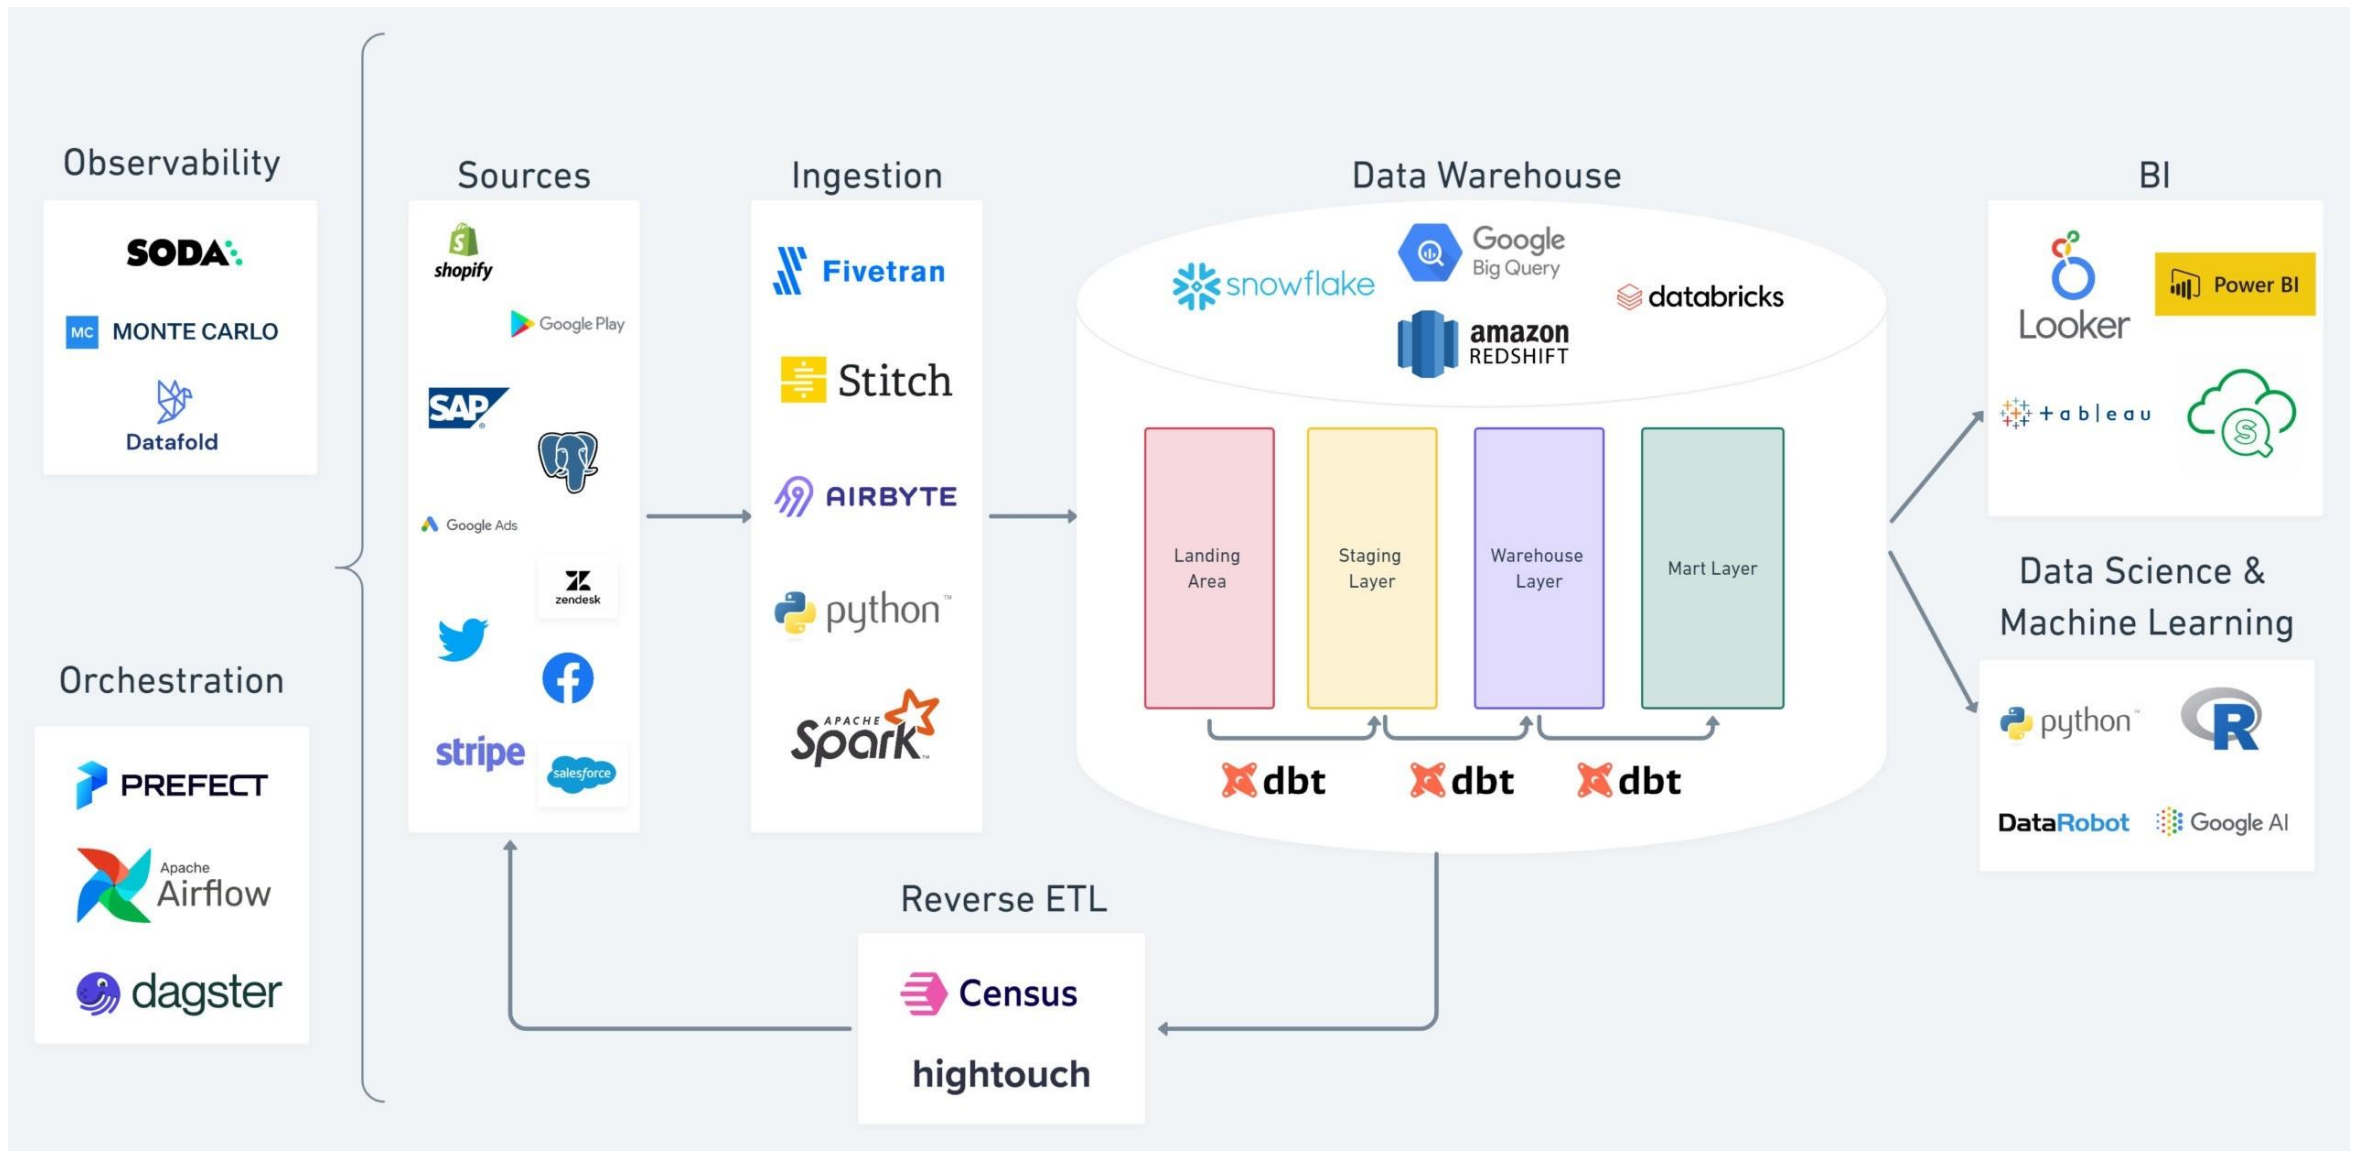
@UNICEFNigeria/2020

Data quality questions?

- Is that table the same compared to the upstream table?
- This table needs to have these columns
- Where this data come from?
- Is the data arrived late?
- Is the values of a numeric column in an expected range?
- Is the summary stats make sense?
- Seem to have outliers?
- We expect the column to have no null value?
- ...

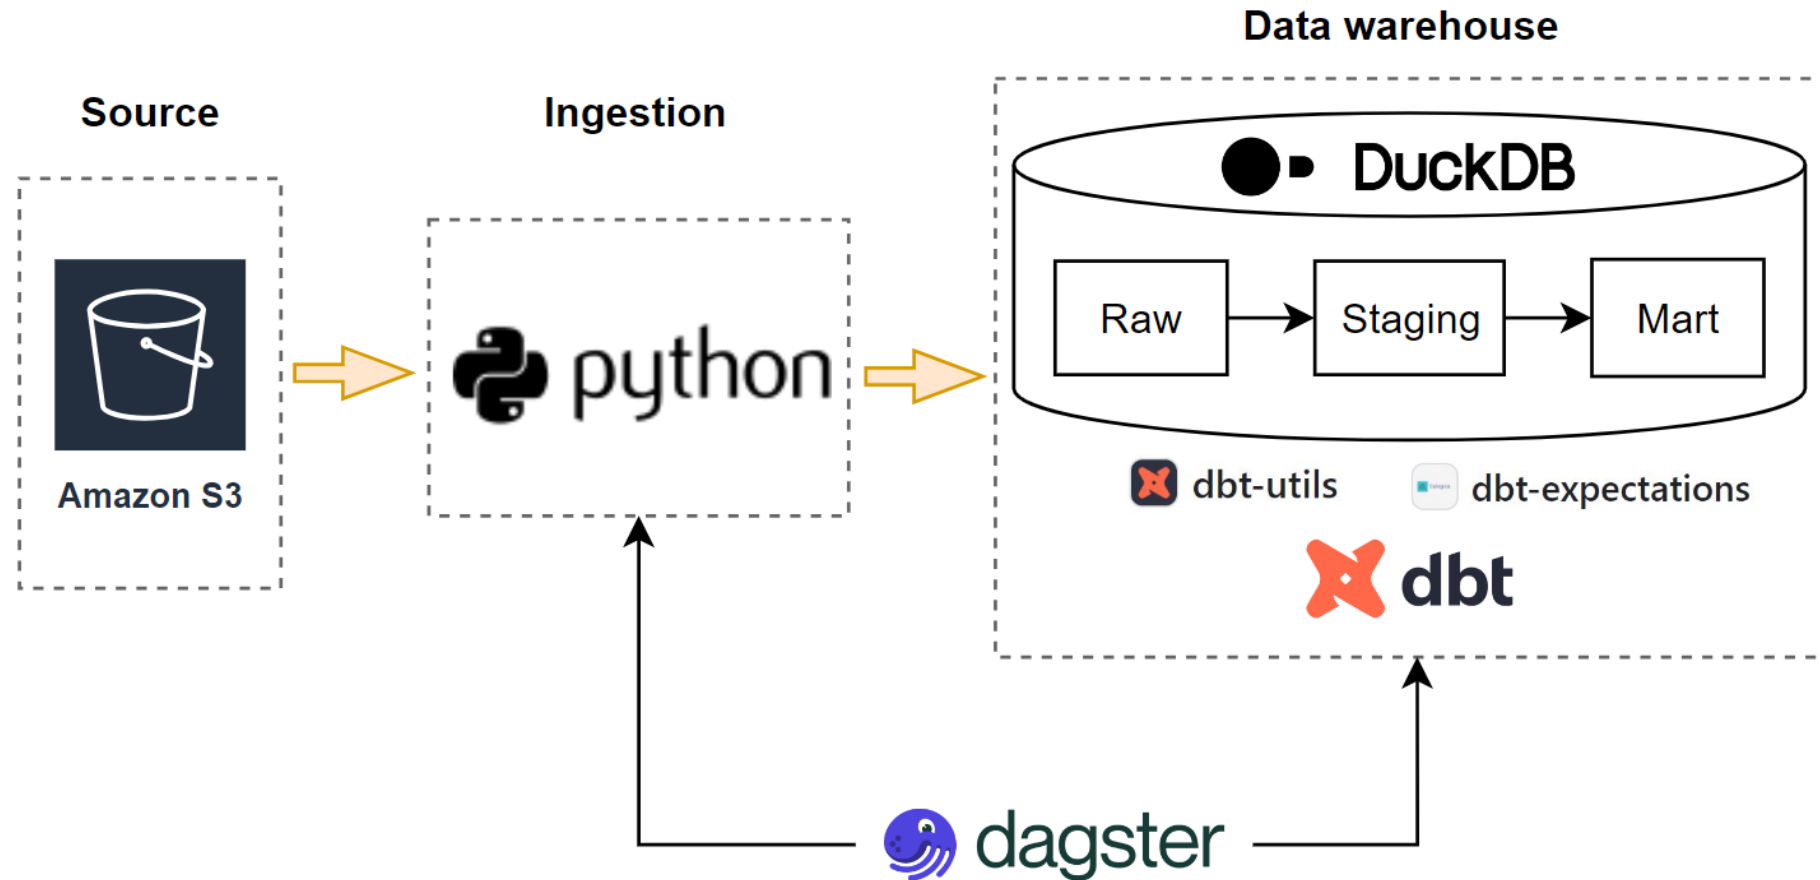
It is critical to test and measure data quality!

**IN DATA
WE TRUST**




https://bitstreams.com/nl/blog/modern_data_stack/

3D (Dagster, Dbt, and Duckdb) trusted pipeline



The 3Ds

Converting an ETL script
to Software-Defined
Assets



DuckDB is an in-process
SQL OLAP database management system

Testing data quality

- **Between tables**

- Relationships
- Raw (unprocessed) truth vs processed truth

- **In a table**

- Row counts
- Column counts
- Column exist
- ...

- **Column level:**

- Generic test:
 - Not null (with a threshold level)
 - Unique
- Text data:
 - Following a pattern?
 - Accepted values?
- Numeric data:
 - Min, max, mean, median
 - Outliers
- Date:
 - Recency, range, latest, min, max

Testing data quality

Four generic tests already defined in Dbt:

- unique,
- not_null,
- accepted_values
- relationships.

```
version: 2

models:
  - name: orders
    columns:
      - name: order_id
        tests:
          - unique
          - not_null
      - name: status
        tests:
          - accepted_values:
              values: ['placed', 'shipped', 'completed', 'returned']
  - name: customer_id
    tests:
      - relationships:
          to: ref('customers')
          field: id
```

Testing data quality

- Add testing packages into dbt *packages.yml* file:

```
1 packages:
2   - package: dbt-labs/dbt_utils
3     version: 1.1.1
4   - package: calogica/dbt_expectations
5     version: [">=0.8.0", "<0.9.0"]
```

- Install dbt packages: **dbt deps**
- Run test: **dbt test**

Testing data quality

- **Dbt-utils** Generic Tests

<https://github.com/dbt-labs/dbt-utils#generic-tests>

- `equal_rowcount` (source)
- `fewer_rows_than` (source)
- `equality` (source)
- `expression_is_true` (source)
- `recency` (source)
- `at_least_one` (source)
- `not_constant` (source)
- `not_empty_string` (source)
- `cardinality_equality` (source)
- `not_null_proportion` (source)
- `not_accepted_values` (source)
- `relationships_where` (source)
- `mutually_exclusive_ranges` (source)
- `sequential_values` (source)
- `unique_combination_of_columns` (source)
- `accepted_range` (source)
- Grouping in tests

Testing data quality

■ Dbt-expectations

<https://github.com/calogica/dbt-expectations>

Table shape

- `expect_column_to_exist`
- `expect_row_values_to_have_recent_data`
- `expect_grouped_row_values_to_have_recent_data`
- `expect_table_aggregation_to_equal_other_table`
- `expect_table_column_count_to_be_between`
- `expect_table_column_count_to_equal_other_table`
- `expect_table_column_count_to_equal`
- `expect_table_columns_to_not_contain_set`
- `expect_table_columns_to_contain_set`
- `expect_table_columns_to_match_ordered_list`
- `expect_table_columns_to_match_set`
- `expect_table_row_count_to_be_between`
- `expect_table_row_count_to_equal_other_table`
- `expect_table_row_count_to_equal_other_table_times_factor`

Missing values, unique values, and types

- `expect_column_values_to_be_null`
- `expect_column_values_to_not_be_null`
- `expect_column_values_to_be_unique`
- `expect_column_values_to_be_of_type`
- `expect_column_values_to_be_in_type_list`
- `expect_column_values_to_have_consistent_casing`

Sets and ranges

- `expect_column_values_to_be_in_set`
- `expect_column_values_to_not_be_in_set`
- `expect_column_values_to_be_between`
- `expect_column_values_to_be_decreasing`
- `expect_column_values_to_be_increasing`

Testing data quality

▪ Dbt-expectations

<https://github.com/calogica/dbt-expectations>

Multi-column

- `expect_column_pair_values_A_to_be_greater_than_B`
- `expect_column_pair_values_to_be_equal`
- `expect_column_pair_values_to_be_in_set`
- `expect_compound_columns_to_be_unique`
- `expect_multicolumn_sum_to_equal`
- `expect_select_column_values_to_be_unique_within_record`

Distributional functions

- `expect_column_values_to_be_within_n_moving_stdevs`
- `expect_column_values_to_be_within_n_stdevs`
- `expect_row_values_to_have_data_for_every_n_datepart`

String matching

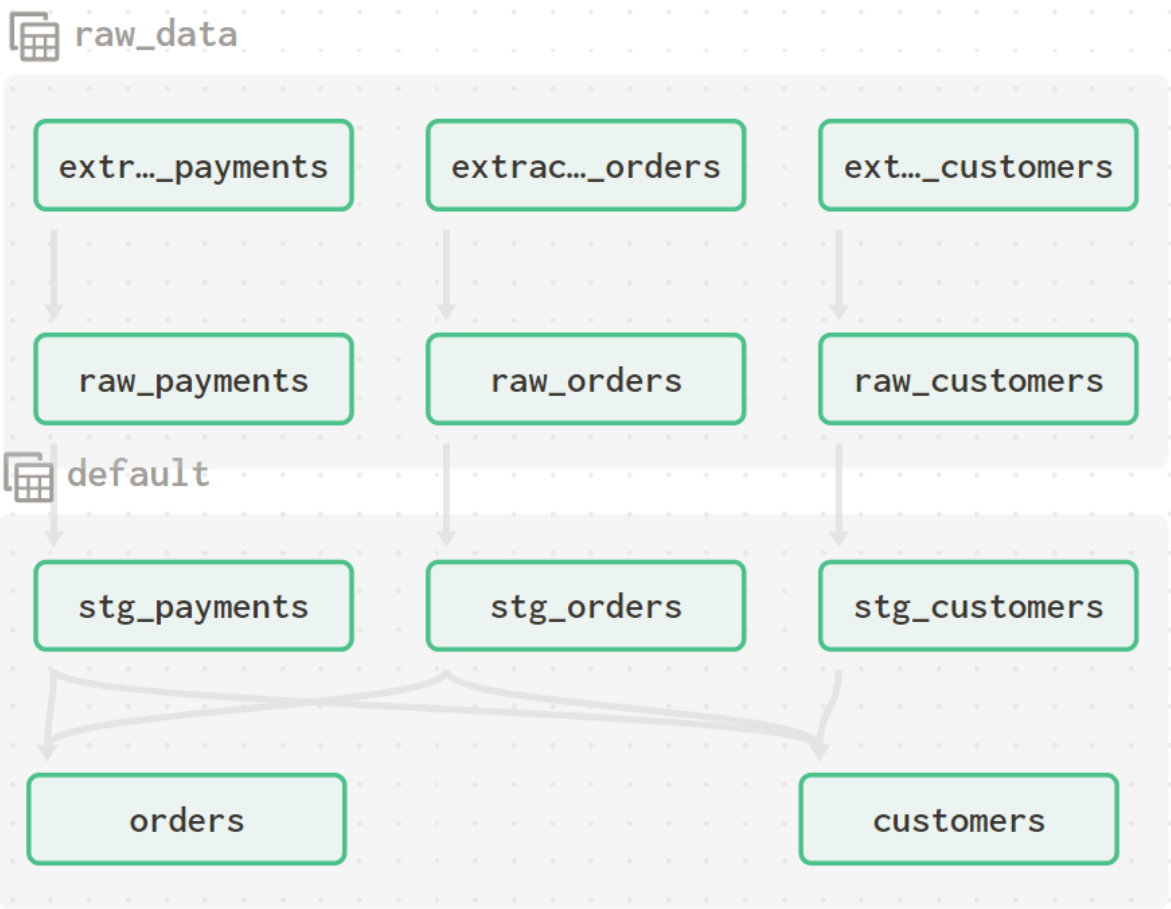
- `expect_column_value_lengths_to_be_between`
- `expect_column_value_lengths_to_equal`
- `expect_column_values_to_match_like_pattern`
- `expect_column_values_to_match_like_pattern_list`
- `expect_column_values_to_match_regex`
- `expect_column_values_to_match_regex_list`
- `expect_column_values_to_not_match_like_pattern`
- `expect_column_values_to_not_match_like_pattern_list`
- `expect_column_values_to_not_match_regex`
- `expect_column_values_to_not_match_regex_list`

Global Asset Lineage

↻ Reload definitions

✳ Type an asset subset... (ex: ++jaffle_)

0:15 ↻ ✳ Materialize all ▼



Search, zoom, and download controls.

Global Asset Lineage

↻ Reload definitions

✂ Type an asset subset... (ex: ++jaffle_

0:00 ↻

⚡ Materialize all ▾

📁 raw_data

📄 extract_raw_payments

No description

Materialized 30 July, 5:48 pm

📄 extract_raw_orders

No description

Materialized 30 July, 5:48 pm

📄 extract_raw_customers

No description

Materialized 30 July, 5:48 pm

📄 raw_payments

No description

Materialized 30 July, 5:49 pm

📄 raw_orders

No description

Materialized 30 July, 5:49 pm

📄 raw_customers

No description

Materialized 30 July, 5:49 pm

📁 default

🔍

🔍

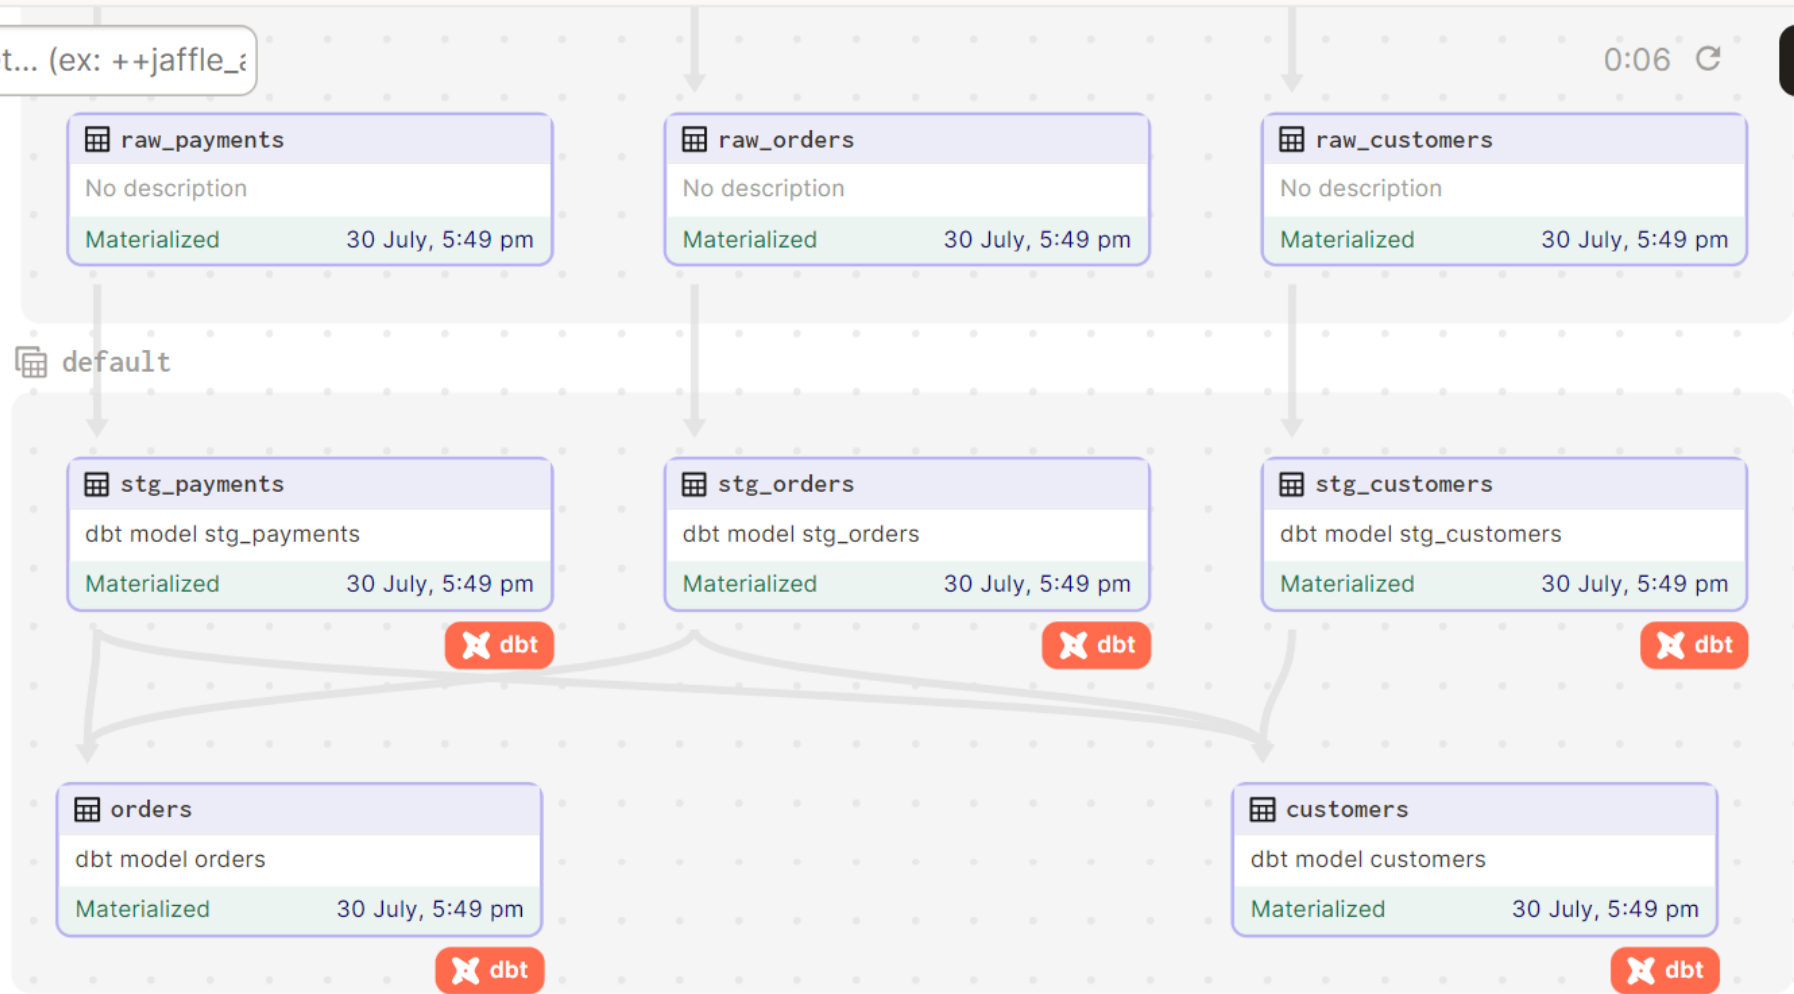
Global Asset Lineage

🔄 Reload definitions

✂ Type an asset subset... (ex: ++jaffle_

0:06 ↺

✨ Materialize all ▾



Vertical toolbar with icons for zoom in (+), zoom out (-), search (magnifying glass), and download (download icon).

Demo

Building trust is a journey!

- Building trusted data is a journey, which needs a good road map
- And strong leaderships with a right data strategy
- And buy-in of data stakeholders and data customers
- Well designed data schema
- Control data quality from upstream (data contract), on the way (data warehouse / data lake), and downstream

References

- The Demo Project repo : <https://github.com/danhphan/trusted-data-pipeline>
- Dbt tests: <https://docs.getdbt.com/docs/build/tests>
- Dbt Workshop: Advanced Testing: <https://www.youtube.com/watch?v=fo7lUn6vgtg>
- Build a poor man's data lake from scratch with DuckDB : <https://dagster.io/blog/duckdb-data-lake>
- Building a robust data pipeline with Dbt, Airflow, and Great Expectations: <https://www.getdbt.com/coalesce-2020/building-a-robust-data-pipeline-with-dbt-airflow-and-great-expectations/>
- Dagster repo: <https://github.com/dagster-io/dagster>
- Dbt repo: <https://github.com/dbt-labs/dbt-core>
- Duckdb repo: <https://github.com/duckdb/duckdb>
- Dbt-utils repo: <https://github.com/dbt-labs/dbt-utils>
- Dbt-expectation repo: <https://github.com/calogica/dbt-expectations>

Thank you!